

UNIVERSITY COLLEGE LONDON  
DEPARTMENT OF COMPUTER SCIENCE  
COMP0023 Networked Systems  
Individual Coursework 2: Implementing Distance-Vector Routing  
**Due: 10th December 2019, 2:55 PM**

In this coursework, you will write distance-vector routing code for a simple router. A valuable reference for you to use while working on this coursework is the set of COMP0023 lecture notes on distance-vector routing. Those slides contain a complete statement of the distance-vector routing algorithm in pseudocode, and examples of how the algorithm behaves on a variety of network topologies. You will also find it useful to refer to RFC2453 (RIP Version 2), which describes the Routing Information Protocol (RIP), a simple distance-vector routing protocol. Note that we expect you to implement the algorithm as presented in the lecture notes, including setting of the appropriate routing table entries' metrics to a reserved INFINITY value when a link goes down. We will also ask you to implement split horizon with poison-reverse, and timeout-based expiration of routing table entries. Note that you are *not* to implement other “advanced” features of DV routing, such as triggered updates.

You must write your code in Java, which is the language we used for the network simulator we give you as a starting point. The code you will submit will be evaluated under Linux on the CS department's lab machines. We have ensured that the code we give you to use as a starting point works correctly on these lab machines. Note that these machines are accessible over the Internet, so you may work on the coursework either from home or in the labs.<sup>1</sup>

## A Simple Network Simulator

A router isn't much good unless it's connected to other routers by links. If you were writing routing protocol software for a real, physical IP router, you could test the software by connecting several routers into a network topology, running your routing software on each of them, and observing whether users' data packets reach their destinations successfully.

Because it's not feasible for you to build a physical multi-router network to test your routing code, we'll do the next-best thing: we'll use “virtual” routers rather than physical ones. We provide you with code for a simple *network simulator* that models a set of routers connected together by a set of links. I.e., the simulator reads a configuration file that describes a network topology, and it then simulates that network by running one copy of your routing code on each router, and passing packets between routers over the links listed in the configuration file.

In this coursework, there is a rigidly defined interface for the router's code. In the simulator's configuration file, you specify not only the network topology, but the name of the compiled Java module for the router code you would like to run on each router in the network. In this way, you can actually mix different router implementations in a single simulated network!

We have provided you with a skeleton of the code for a distance-vector router, found in the file `DV.java`. You should implement your solution to the coursework by filling in the missing parts of this file. Do not change any of the constants that we've pre-defined in that file—they must be left as-is for the simulator to work properly. The skeleton adheres to the `RoutingAlgorithm` interface to the distance-vector router, which you *must not change*; any

---

<sup>1</sup>Because of the particular configuration of the CS department's network, if you would like to use any of the lab machines remotely, you must first log into a CS departmental gateway such as `newgate.cs.ucl.ac.uk` by `ssh`, then from there log into one of the lab machines using `ssh` again.

changes to the router's API will similarly cause the simulator not to work correctly! The skeleton code compiles, but all methods in it return dummy or null return values. To build your solution, you must implement the following in the file `DV.java`:

- the bodies of the methods in the `DV` class that are incomplete (return dummy or null return values) in the version of `DV.java` we've given you. The interface `RoutingAlgorithm` in `RoutingAlgorithm.java` includes the description of what to implement in every `DV` method.
- all methods in `DVRoutingTableEntry`, a class that implements the `RoutingTableEntry` interface found in `RoutingTableEntry.java`

Note that you should *not modify any files* in the coursework apart from `DV.java` and the configuration files for the simulator, described below: all the code you write will go in `DV.java`.

We have given you a `Makefile` (found in the set of files for the coursework) to help automate the compiling and running of your routing code and the simulator. A full description of the `make` utility is beyond the scope of this coursework. For the purposes of this coursework, all you need to know about compiling and running your code is:

- Don't modify the `Makefile`.
- To compile your routing code in `DV.java` into the compiled Java module `DV.class`, just type

```
make
```

in the same directory where the `Makefile` and all the source files are located.

- To prepare the documentation, just type `make javadoc` in the directory containing your coursework files, and you will find the documentation for the code we've given you in a newly created `docs` sub-directory.
- To run tests of your routing protocol implementation, after you've compiled your router's code with `make` as above, just type

```
java Simulator config.cfg
```

where `config.cfg` can be the name of any simulator configuration file (whose format we describe below).

- To see brief help on what functions the `Makefile` lets you automate, type:

```
make help
```

## Simulator Configuration File

Each time you run the simulator, it reads a configuration file that describes the particular network topology it should simulate, and any actions to take during the simulation (and when to take them), such as “take this link down after 15 seconds,” “print out the routing table of this router after 32 seconds,” *etc.* – as described further below.

Consider the following simple example configuration file:

```
updateInt 10

preverse off
expire off

router 0 2 DV
router 1 2 DV
router 2 2 DV

link 0.0.1 1.0.1
link 1.1.1 2.0.1
link 2.1.1 0.1.1

send 10 0 1

downlink 10 1.1 2.0

uplink 12 1.1 2.0

dumpPacketStats 14 all

dumprt 14 all

stop 100
```

The first three non-empty lines of the above configuration file specify that all routers in the network should send DV protocol updates every 10 seconds, that split horizon with poison-reverse should be disabled at all routers, and that timeout-based expiration of routing table entries should also be disabled at all routers.

Thereafter, the configuration file describes a scenario involving three simulated routers with addresses 0, 1, and 2, arranged in a ring. 10 seconds into the simulation, router 0 originates a data packet (to be forwarded by the routers in the network) with destination address 1. Also 10 seconds into the simulation, the link between router 1 and router 2 goes down. This link comes back up 12 seconds into the simulation.

All routers dump summary statistics of how many packets they’ve sent, received, dropped, and forwarded after 14 seconds, and dump their routing tables after 14 seconds.

The simulation runs for 100 seconds.

Router IDs are simple integers, as are interface IDs on routers.

Now, let’s fully define the syntax of lines in the configuration file.

The routing algorithm update interval is declared as follows:

**updateInt** *u*

where *u* is the interval between DV protocol updates for all routers defined thereafter in the configuration file. If **updateInt** is not declared in the configuration file, then all routers use a default update interval of 1 second.

One may also enable or disable specific DV protocol optimizations at all routers. Enable or disable split horizon with poison-reverse as follows:

**preverse** on | off

where *on* or *off* specifies whether or not the routing algorithm deployed on each router incorporates split horizon with poison reverse.

Similarly, one may enable or disable timeout-based expiration of table entries as follows:

**expire** on | off

where *on* or *off* specifies whether or not the routing algorithm deployed on each router should expire routing entries in its routing table based on a timeout interval.

A router is declared as follows:

**router** *id n classname*

where *id* is the integer ID of this router, *n* is the number of interfaces for the router, *classname* is the name of the compiled Java module that should be used for the routing software for this router.

A link is a connection between two routers. Links also have a metric in each direction (configured in real routers by the network administrator). Links are declared as follows:

**link** *rlid.rlif.rlw r2id.r2if.r2w* [up | down]

where *rlid* is the integer ID of the router at one link endpoint, *rlif* is the interface ID to which the link connects on *rlid*, *rlw* is the metric incurred by packets sent by *rl* on the link, and all the *r2\** fields have the same meanings for the router at the other link endpoint. The last field in the **link** line is optional. If supplied, it defines the initial state of the link in the simulation as either up or down.

You control the length of the simulation with:

**stop** *time*

where *time* is the number of seconds to run the simulation (the clock starts at zero seconds).

To allow you to observe how packets are routed in the network, the simulator allows injecting data packets from a particular source to a particular destination. You can do so with:

**send** *time origin destination*

where *time* is the number of seconds into the simulation to originate the packet, *origin* is the ID of the router to send the packet, and *destination* is the destination ID to put into the packet.

To see how the routing system behaves when links go down and come back up, the simulator supports taking links down and up at specified times. To take a link down, use:

```
downlink time router1.interface1 router2.interface2
```

where *time* is the time to break the link, *router1* and *interface1* are the router ID and interface ID of one end of the link, and *router2* and *interface2* are the router ID and interface ID of the other end of the link. Note that the first router specified should be the one with the smaller router ID.

Similarly, you can bring a link up with:

```
uplink time router1.interface1 router2.interface2
```

where the parameters are the same as those for `downlink`.

Finally, the simulator supports two commands to let you inspect the internal state of routers at a specified point in time. To see how many packets have been sent (s), received (r), dropped (d), and forwarded (f) by a router, use:

```
dumpPacketStats time router|all
```

where *time* states when you'd like packet statistics, and either *router* specifies the ID of a single router where you'd like packet statistics, or *all* specifies that you'd like packet statistics from all routers.

To see a router's routing table, use:

```
dumpprt time router|all
```

where the parameters are the same as those for `dumpPacketStats`. (Note that the function that outputs a router's current routing table is unimplemented in the initial code we've given you. You must implement this functionality; we define the format in which you must output routing tables in the next section.)

The commands `dumpprt` and `dumpPacketStats` will be very useful to you in debugging your router—if your router doesn't behave as you expect it to, you can add these commands to a configuration file to view the routing tables and packet statistics at any step in time you like. Note that those commands are executed at the beginning of the simulator's steps, before packets are processed during that step – see, e.g., the `main_loop()` function in `Simulator.java`.

## Completing the Coursework

The first step in getting started with the coursework is to make yourself a copy of the files we give you to start from. Download the archive for the coursework (`dv-cw.tar.gz`) from the course Web page and execute the following command:

```
tar xvzf dv-cw.tar.gz
```

You will then find a new directory `dv-cw` in your current directory, which contains all the coursework files.

When we mark your coursework, we will use a series of tests for your router. In each test, we will run the simulator with your routing code on every router, using a configuration file with a test network topology.

There are several topologies on which we'll test your router. We've given you five of them: these are in the simulator configuration files named `test1.cfg` through `test5.cfg`. We describe these test cases in more detail further below. There are additional test cases with which we'll test your router, too, known only to the course staff. We hold these in reserve until marking time so that you have an incentive to make sure your router truly works correctly for all topologies, rather than trying to "target" your implementation to the five tests we've given you.

There are three stages in which you should complete this coursework. In each stage, there is functionality you must implement, and the functionality is cumulative across the three stages.

## Stage 1: Baseline DV

In this stage, you must implement a baseline DV routing algorithm. No separate design document is required, but you must comment your code thoroughly, to fully explain how it works.

The first two tests, `test1.cfg` and `test2.cfg`, check the correctness of your baseline DV router implementation. They therefore disable split horizon with poison-reverse and timeout-based table entry expiration. *Leave these two features disabled in these two tests' configuration files! You will be marked based on your router's behavior on these two tests with these two features disabled.*

## Stage 2: Add Split Horizon with Poison-Reverse

In this stage, you must add split horizon with poison-reverse (SH/PR) to improve the convergence behavior of your DV routing implementation. We've given you two test cases, `test3.cfg` and `test4.cfg`, that have `preverse` set to on. That's the configuration in which we will use these two tests when marking your submission, but the first part of this stage is to explore how your baseline router *without* SH/PR behaves on these topologies.

**Written question (answer to be turned in):** Without SH/PR enabled, run `test3.cfg` and `test4.cfg`. Consider the link failures that occur during these simulations. When links fail, what routing and forwarding pathologies, if any, do you observe in each of these two test configurations when DV does not use SH/PR?

Now implement SH/PR. Note that you will need to enable or disable SH/PR in your router in accordance with the setting of the `preverse` flag in the configuration file. The code we've given you reads that statement from the configuration file already; read the Javadoc documentation to see which variables to access in the DV implementation to obtain the value of this flag. (When testing your SH/PR implementation, be sure the `preverse` flag is on! And similarly, if you want to run without SH/PR after implementing them, be sure to turn this flag off.)

**Written question (answer to be turned in):** Now with SH/PR enabled, examine the behavior of your DV implementation on `test3.cfg` and `test4.cfg`. For each of the two tests, does SH/PR prevent the pathologies you previously observed? Explain why or why not.

(The submission instructions later in this document describe how to submit your answers to these questions with your code.)

## Stage 3: Add Expiration of Stale Table Entries

In this stage, you must enhance your DV implementation further to expire routing table entries as necessary. Note that the DV algorithm described in lecture allows routing table entries to persist indefinitely. But if a destination goes down and stays down, routers do not need to maintain routing table entries for that destination. In fact, deleting such entries from routing tables will reduce the size of subsequently exchanged routing messages – although it doesn’t hasten convergence.

You will need to enforce a deadline after which stale entries for unreachable destinations should be removed from the routing table. We require that you remove stale entries in accordance with the “garbage-collection” timeout policy specified in RFC2453 (which you will find a very useful reference). In particular, you should keep the same proportion between the “garbage-collection” timeout and the update interval – i.e., have the timeout four times bigger than the update interval. Note that a second “timeout” is specified in the RFC, mostly to handle the loss of routing update packets. Your simulator, however, never drops routing update packets, so you do not need to incorporate the second “timeout” or similar mechanisms for dealing with update losses.

`test5.cfg` tests for the removal of stale routing table entries. This test enables both SH/PR and stale table entry expiration—the configuration in which we will test your code on `test5.cfg` during marking.

## Further Tips on Implementation and Testing

### Getting Marks

To pass automated testing of your router, you **must** adhere to the following format when you implement the `dumprt` command:

- Output, on one line:

`Router n`

where you replace *n* with the integer address of the router.

- For each destination in the router’s routing table, print out one line, in the format:

`d destination_id i interface_id m metric`

where you replace *destination\_id* with the integer address that is the destination for this entry, *interface\_id* with the integer interface ID for this entry, and *metric* with the integer metric for this entry.

**Your router should not output any text other than the above as part of its `dumprt` output.** We use automated scripts to test your router’s behavior, and these scripts rely on your strict adherence to the above output format.

You will be marked on whether you pass the five public router tests and course staff’s private router tests, and on the clarity of your design and comments.

## Comparing with a Correct Solution

We enclosed the output of a correct DV implementation within the `dv-cw.tar.gz` file distributed with the coursework. After unpacking the tar file, you will find the output that we expect your implementation to produce for every given test file in a directory `ModelTestOutput` inside the `dv-cw` one.

## Marking Scheme

Out of 100 marks in total for the coursework, we will allocate marks as follows:

- passing `test1.cfg`: 5 marks
- passing `test2.cfg`: 10 marks
- passing `test3.cfg`: 15 marks
- passing `test4.cfg`: 15 marks
- passing `test5.cfg`: 15 marks
- passing two unseen tests: 15 marks
- your written answers to the two questions: 15 marks
- coding style, comment clarity and completeness: 10 marks

## Testing Your Code

You can use `make` to run the tests. Do so by typing:

```
make test1
make test2
make test3
make test4
make test5
```

to run each of the tests in the five test simulator configuration files, respectively. These `make` commands will store the output of the simulator in files named `testNOutput.txt`, where `N` is the number of the test in question.

## What to Turn In, and How

Submission of the coursework is electronic using the Moodle web site. You must turn in both of the following files electronically:

- A **single Java** source code file, `DV.java`, containing your full implementations of *both* class `DV` and class `DVRoutingTableEntry`.



- A **single text** file named `answers.txt` containing your answers to the two questions in Stage 2. In the text file, enumerate (e.g., in a list or table) the downlink events simulated in test3 and test4: for each event, describe the consequences of that event on routing convergence and packet forwarding.

## Academic Honesty

You are permitted to discuss the concepts of distance-vector routing (that is, the lectures' and assigned readings' content) with your classmates, and to discuss debugging strategies with one another, but *you are not permitted to show your code to any other student, in whole or in part, nor are you permitted to contribute any lines of code to any other student's solution*. As one always does in an academic setting, you must acknowledge the work of others explicitly. In this case, that means that if a classmate discussed distance-vector routing with you, or helped you strategize on how to debug your code, you must state the identity of that classmate in what you hand in, and describe how they contributed to your work (clearly indicated in a comment at the top of your Java source code file `DV.java`).

**All code that you submit must be written entirely by you alone.**

We use sophisticated copying detection software that exhaustively compares code submitted by all students from this year's class and past years' classes, and produces color-coded copies of students' submissions, showing exactly which parts of pairs of submissions are highly similar.

**Do not copy code from anyone, either in the current year, or from a past year of the class.** You **will** be caught, just as students have been caught in years past.

Copying of code from student to student is a serious infraction; it will result in automatic awarding of zero marks to all students involved, and is viewed by the UCL administration as cheating under the regulations concerning Examination Irregularities (normally resulting in exclusion from all further examinations at UCL). You have been warned!

## Questions and Piazza Site

If you have questions about the coursework, please don't hesitate to visit us during office hours, or to ask questions on Piazza. When asking questions on Piazza, please be careful to mark your question as private if it reveals details of your solution. (Questions that don't reveal details of your solution, such as those about how to interpret the coursework text or lecture material, should be left public, though, so that all in the class may benefit from seeing the answers.)

As always, please monitor the Piazza site of the course. Any announcements (e.g., helpful tips on how to work around unexpected problems encountered by others) will be posted there.